




cluster configurations

These tests were run on the Converge data (across 10640 samples) on the ODA cluster.






Local mode, Spark 2.1.1, Scala 2.11.8, 48 cores, 32G Driver, materializing shredded query stages


```
--master local[*] \  
--executor-memory 24G \  
--driver-memory 32G \
```

Query4




Converge, 2 slaves, Spark 2.1.1, Scala 2.11.8, 4 partitions, materializing shredded query stages

```
--master spark://192.168.11.235:7077 \  
--executor-cores 4 \  
--num-executors 8 \  
--executor-memory 24G \  
--driver-memory 32G \
```



Converge, 2 slaves, Spark 2.1.1, Scala 2.11.8, 8 partitions, materializing shredded query stages

```
--master spark://192.168.11.235:7077 \  
--executor-cores 4 \  
--num-executors 8 \  
--executor-memory 24G \  
--driver-memory 32G \
```



Converge, 1 slave, Spark 2.1.1, Scala 2.11.8, 4 partitions, materializing shredded query stages

```
--master spark://192.168.11.235:7077 \  
--executor-cores 4 \  
--num-executors 4 \  
--executor-memory 24G \  
--driver-memory 32G \
```

